

# Image Classification Using CNN from Optimization Perspective

Farida Far Poor, Mohammad Mehdi Hosseini



# Table of contents

1. Classification
2. Neural Network
3. Why CNN
4. Our Network Structure
5. Optimization Perspective
6. Experiments and Analysis



# Classification

## Classification:

- Design a model
- Look at many samples (providing data)
- Learn features (loss minimization)
- Label new **unseen** data

## Classifier is


- a **model** that produces **labels** from a set of **features** of a **data**

## Image classifier:

- Model
- Utilize Data (images)
- Train itself
- Label data

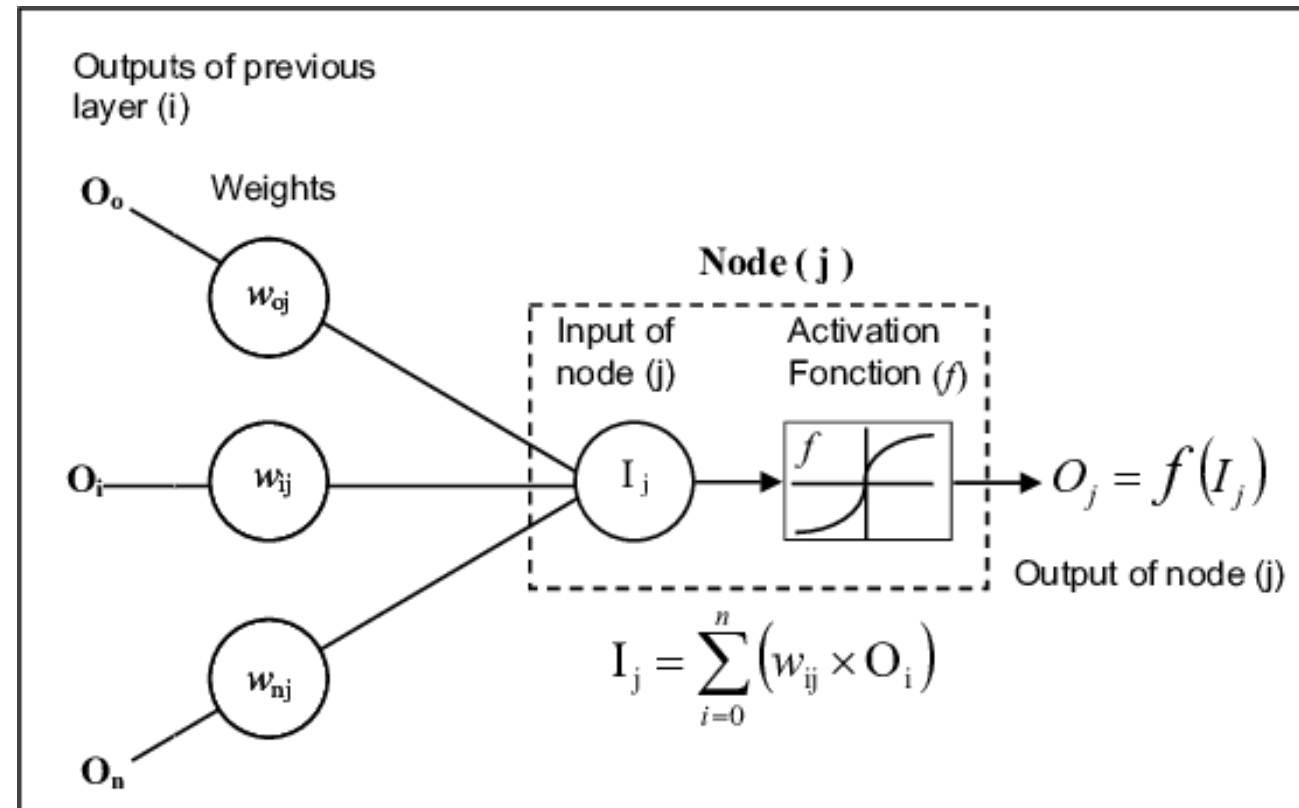
# Classic Neural Network

## Feed forward + Back-propagation



# Classic Neural Network

## Inside a node



# Traditional Methods Drawbacks

Main problem in classification:

Feature extraction is

**boring**

**time consuming**

**not exact**

Features like: HOG, SIFT, HSV, RGB are understandable by **human**, not necessarily by the **computers**.



# Classic Neural Network

Other assistants:

Feature selection

Feature generation

Painkillers



# Question

Computers' understanding is not quite the same as humans,

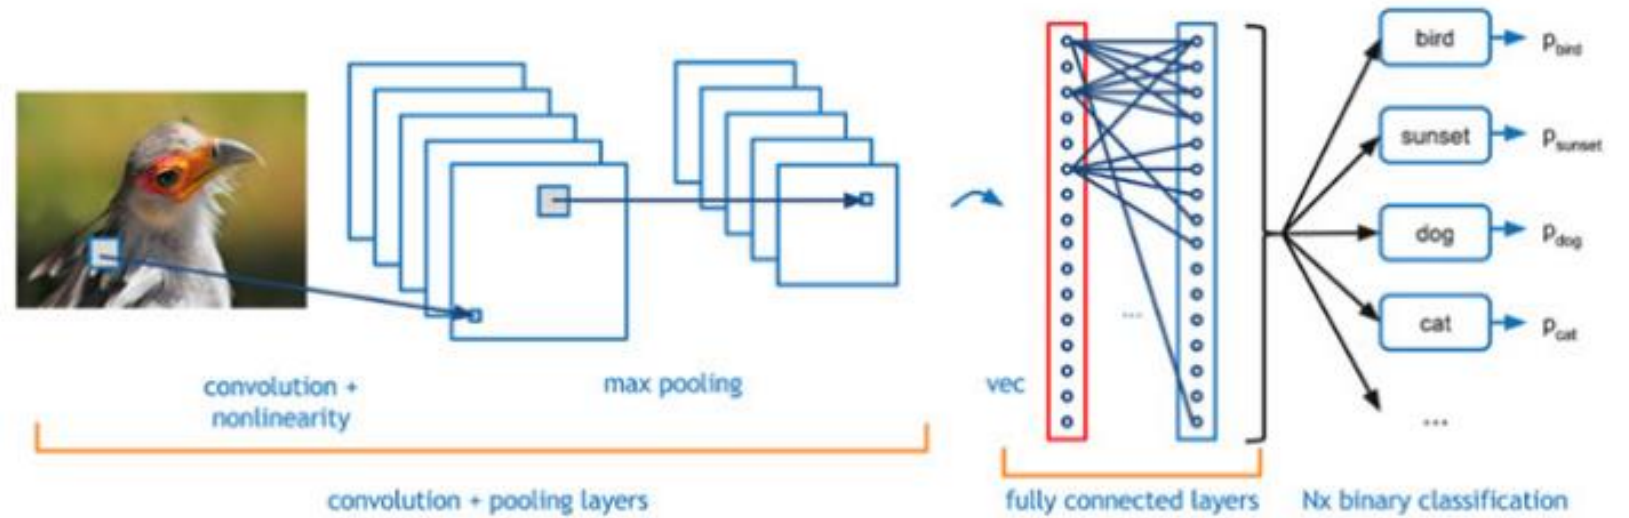
Isn't it better **let them do that?**



Now, we know why we use convolutional neural networks

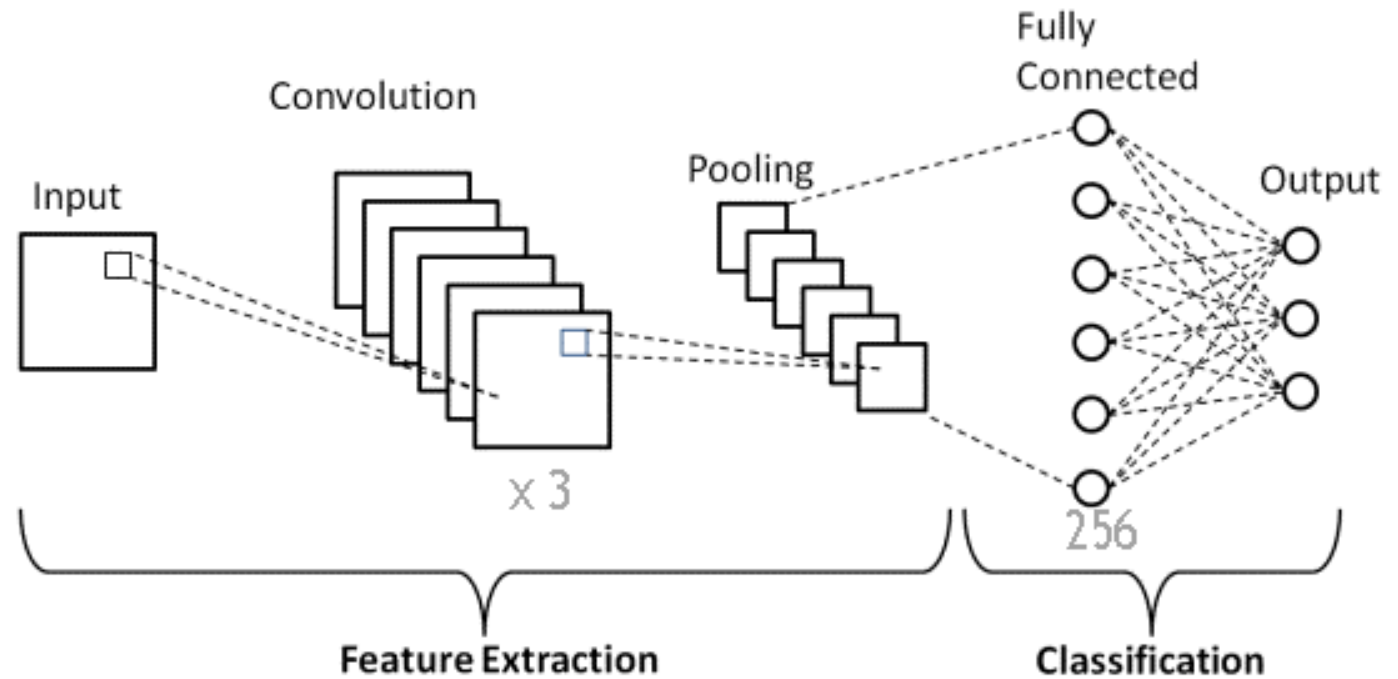


# Convolutional Neural Network



# Network Structure

3 convolutional layers  
1 fully connected layer




# Optimization Perspective

## ❖ Maximize Accuracy (Minimize Error)


- Loss Function
- Optimizer
  - Learning Rate
- Activation function

# Loss Function


✓ MSE



✓ Cross-entropy



✓ Hinge



# Mean Squared Error (MSE)

General Form:

$$\theta^* = \arg \min_{\theta} f(\theta)$$

MSE (suitable for regression):

$$f(\theta) = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i(\theta))^2$$

For example, in a 2D space:

$$\min_{w_0, w_1} \frac{1}{n} \sum_{i=1}^n (y_i - (w_0 + w_1 \cdot x_i))^2$$

$w_1 \rightarrow$  slope       $w_0 \rightarrow$  intercept

# Cross-Entropy

General Form:

$$\theta^* = \arg \min_{\theta} \mathcal{L}(\theta)$$

Cross entropy (binary classifier):

$$\text{Loss} = -\frac{1}{n} \sum_{i=1}^n (y_i \cdot \log(\hat{y}_i) + (1 - y_i) \cdot \log(1 - \hat{y}_i))$$

Categorical Cross Entropy (Multi-class classifier):

$$\text{Loss} = -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^C y_{i,j} \cdot \log(\hat{y}_{i,j})$$

CE is a 0/1 classifier.

# Hinge

General Form:

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta})$$

Hinge (binary):

$$\text{Hinge Loss} = \frac{1}{n} \sum_{i=1}^n \max(0, 1 - y_i \cdot f(\mathbf{x}_i))$$

Hinge with constraint:

$$\boldsymbol{\theta}^* = \arg \min_{\boldsymbol{\theta}} \left( \frac{1}{n} \sum_{i=1}^n \max(0, 1 - y_i \cdot f(\mathbf{x}_i; \boldsymbol{\theta})) \right) + \lambda \|\boldsymbol{\theta}\|^2$$

Hinge loss is a -1/1 classifier.

# Optimizers

## Stochastic Gradient Descent (SGD)

$$\theta_{t+1} = \theta_t - \eta \nabla J(\theta_t; x^{(i)}, y^{(i)})$$

$\theta_t$  → Parameters at time t

$\eta$  → Learning rate

$\nabla J(\theta_t; x^{(i)}, y^{(i)})$  → Gradient

Works better using batches.



# Optimizers

ADAM

$$\theta_{t+1} = \theta_t - \eta \frac{\hat{m}_t}{\sqrt{\hat{v}_t + \epsilon}}$$

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$


$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) \nabla J(\theta_t)$$

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) (\nabla J(\theta_t))^2$$


How to choose learning rate?

# Activation Function

Sigmoid



ReLU



# Experiments

- ❑ **Classification on two datasets:**
  - Caltech 101
  - Cifar 100
- ❑ **Study network performance:**
  - Different loss functions
  - Different optimizers
  - Different activation functions
  - The best learning rate



# Baseline

**Dataset: Caltech**

**Train-test ratio: 70-30%**

**Loss: categorical crossentropy**

**Optimizer: adam**

**Learning rate: 0.00003**

**Activation function: ReLU**

**Pooling: max**

**Batch size = 32**

**Epoch: 20**



# Datasets

## 1. Caltech 101

102 classes

9000 images

each class 40-800 images

different image size (200~300)



## 2. Cifar 100

100 classes

60000 images

each class 600 images

32x32




# Caltech

Accuracy on Model

Image Size: 256x256

Parameters ~ 15 M




# CIFAR

Accuracy on Model


Image Size: 32x32

Parameters ~ 250 K



# Learning Rate


Adam optimizer






# Learning Rate (Caltech)

## Adam optimizer



# Learning Rate (Cifar)

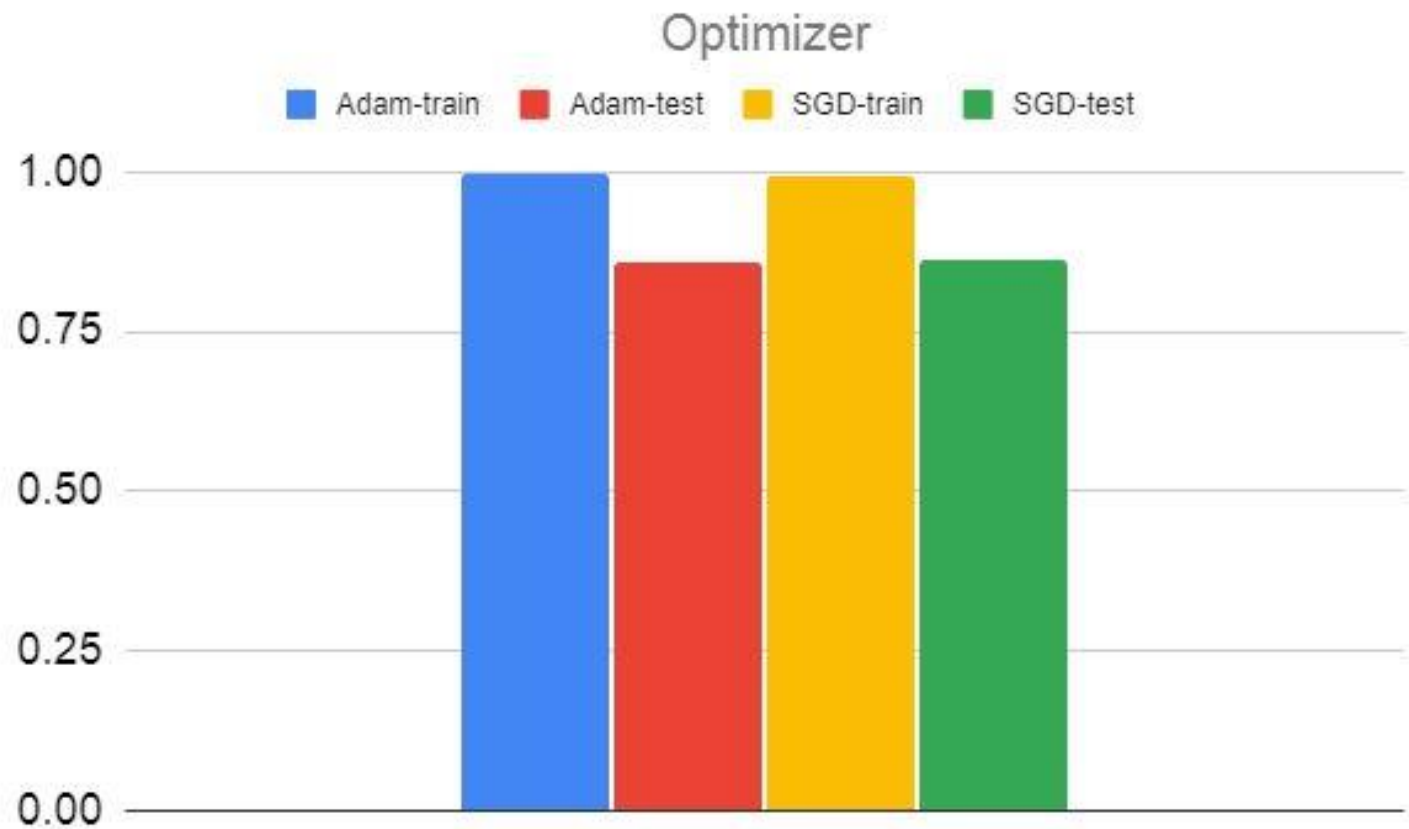
Adam optimizer



# Optimizer (Caltech)

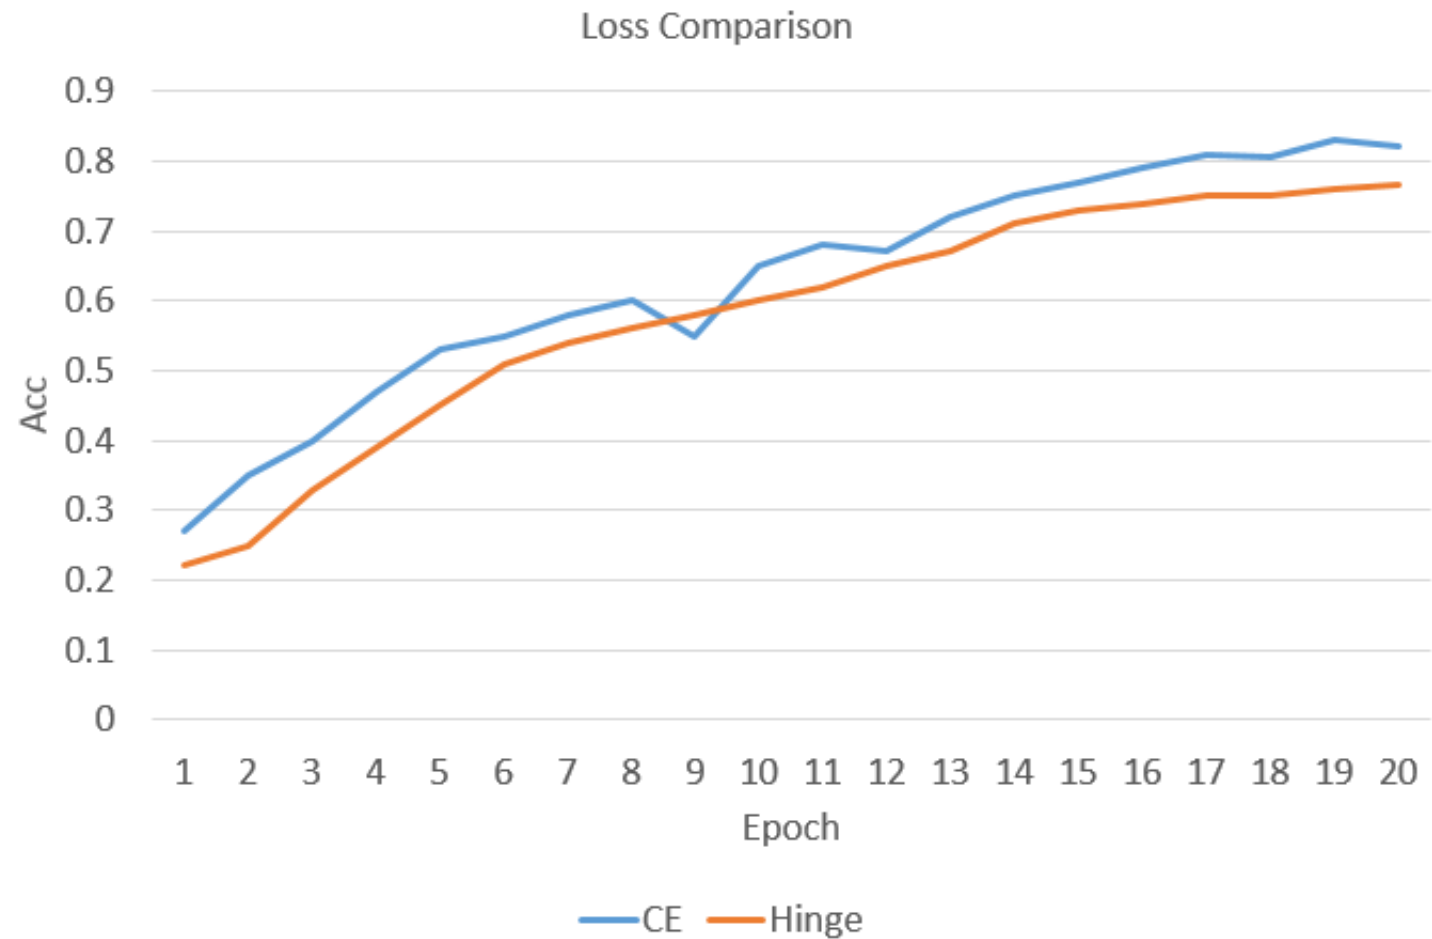
Adam vs SGD

Almost similar results




# Loss Function (Caltech)

## Accuracy using different loss functions



# Activation Function (Caltech)

## Similar results



Question?



<https://openai.com/dall-e-3>